

```
In [1]: #imports need packages
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sn
```

```
In [2]: #Loads the data set and shows the shape of it
data_df=pd.read_csv("C:/Users/predi/Documents/GitHub/DSC630/Datasets/airline_2m.csv", encoding='latin-1')
data_df.shape
```

C:\Users\predi\AppData\Local\Temp\ipykernel\_19772\2826587435.py:2: DtypeWarning: Columns (69,76,77,84) have mixed types. Specify dtype option on import or set low\_memory=False.

```
data_df=pd.read_csv("C:/Users/predi/Documents/GitHub/DSC630/Datasets/airline_2m.csv", encoding='latin-1')
Out[2]: (2000000, 109)
```

## Data Cleaning

```
In [3]: #strips white sapce and lowers cases all of the headers and text.
data_df.columns=data_df.columns.str.lower()
data_df.columns=data_df.columns.str.strip()
```

```
In [4]: #Checks for null values in the data set and shows all columns with missing values
columns_with_null = data_df.columns[data_df.isnull().any()]
columns_null_count = data_df[columns_with_null].isnull().sum()
columns_above_90na = columns_null_count[columns_null_count > 1800000]
```

```
In [5]: #Drops all columns missing more than 90% of their values
data_df.drop(columns=columns_above_90na.index, inplace=True)
data_df.shape
```

```
Out[5]: (2000000, 61)
```

```
In [6]: #Rechecks for nulls now that largest offends have been removed
columns_with_null = data_df.columns[data_df.isnull().any()]
columns_null_count = data_df[columns_with_null].isnull().sum()
columns_null_count
```

```
Out[6]: tail_number      391763
originstate      646
originstatefips  646
originstatename  646
deststate      594
deststatefips    594
deststatename    594
deptime      36005
depdelay      36068
depdelayminutes  36068
depdel15      36068
departuredelaygroups  36068
taxiout      415642
wheelsoff      415677
wheelson      417958
taxiin      417847
arrtime      39551
arrdelay      41078
arrdelayminutes  41078
arrdel15      41078
arrivaldelaygroups  41078
crselapsedtime  281
actualelapsedtime  41052
airtime      419349
carrierdelay    1778197
weatherdelay    1778197
nasdelay        1778197
securitydelay    1778197
lateaircraftdelay  1778197
divairportlandings  1253886
dtype: int64
```

```
In [7]: #dropping all rows missing the dependant variable needed for the model
data_df.dropna(subset=['arrdelay'], inplace=True)
#dropping all rows with missing tail numbers due to these observations
#missing multiple other key data points
data_df.dropna(subset=['tail_number'], inplace=True)
#Rechecks for nulls now that some of the NaN's have been dropped
columns_with_null = data_df.columns[data_df.isnull().any()]
columns_null_count = data_df[columns_with_null].isnull().sum()
columns_null_count
```

```
Out[7]: depdelay      60
depdelayminutes    60
depdel15          60
departuredelaygroups  60
wheelsoff         34
wheelson         111
airtime           1
carrierdelay      1358816
weatherdelay      1358816
nasdelay          1358816
securitydelay     1358816
lateaircraftdelay  1358816
divairportlandings  849361
dtype: int64
```

```
In [8]: #Drops the NaN value rows for the columns with less than 60.
data_df.dropna(subset=[
    'arrdelay', 'wheelson', 'depdelay', 'wheelsoff', 'airtime'], inplace=True)
#Rechecks for nulls now that the tail number NaN's have been dropped
columns_with_null = data_df.columns[data_df.isnull().any()]
columns_null_count = data_df[columns_with_null].isnull().sum()
columns_null_count
```

```
Out[8]: carrierdelay    1358640
weatherdelay    1358640
nasdelay        1358640
securitydelay    1358640
lateaircraftdelay  1358640
divairportlandings  849216
dtype: int64
```

```
In [9]: #fills the NaN values for the remaining columns that are needed for the model with 0
data_df.fillna(value=0,inplace=True)
#Checks for any other NaN values in the dataframe
data_df.columns[data_df.isnull().any()]
```

```
Out[9]: Index([], dtype='object')
```

```
In [10]: #Dropping string columns that have a string and a matching ID column
redundant_columns=['reporting_airline','iata_code_reporting_airline','origin',
                  'origincityname','originstate','originstatename','dest',
                  'destcityname','deststate','deststatename','distancegroup',
                  'flightdate']
data_df.drop(columns=redundant_columns, inplace=True)
#dropping columns that aren't useful(59 min block every plane has sceduled)
data_df.drop(columns=['deptimeblk','arrtimeblk'], inplace=True)
```

```
In [11]: #Now dropping any columns that would not be known before take off
#This is to make the model actually useful after take off the delay is
#relatively easy to calulate.
future_columns=['deptime','depdelay','depdelayminutes','depdel15','departuredelaygroups',
               'taxiout','wheelsoff','wheelson','taxiin','arrtime','arrdelayminutes',
               'arrdel15','arrivaldelaygroups','cancelled','diverted','actualelapsedtime',
               'airtime','divairportlandings']
data_df.drop(columns=future_columns, inplace=True)
#Shows the final shape of the data frame
data_df.shape
```

```
Out[11]: (1580411, 29)
```

```
In [12]: #Changes the delay times to zeros and ones
#The delay can be expected but the time would not be known until takeoff
data_df['carrierdelay'] = data_df['carrierdelay'].apply(lambda x: 1 if x > 0 else x)
data_df['weatherdelay'] = data_df['weatherdelay'].apply(lambda x: 1 if x > 0 else x)
data_df['nasdelay'] = data_df['nasdelay'].apply(lambda x: 1 if x > 0 else x)
data_df['securitydelay'] = data_df['securitydelay'].apply(lambda x: 1 if x > 0 else x)
data_df['lateaircraftdelay'] = data_df['lateaircraftdelay'].apply(lambda x: 1 if x > 0 else x)
```

```
In [13]: data_df
```

```
Out[13]:
```

	year	quarter	month	dayofmonth	dayofweek	dot_id	reporting_airline	tail_number	flight_number	reporting_airline	o
0	1998	1	1	2	5		19386	N297US		675	
1	2009	2	5	28	4		20437	N946AT		671	
2	2013	2	6	29	6		20398	N665MQ		3297	
3	2010	3	8	31	2		19790	N6705Y		1806	
4	2006	1	1	15	7		20355	N504AU		465	
...	...	...	...	...	...		...	...		...	
1999995	2008	1	3	23	7		19393	N712SW		966	
1999996	1999	1	1	5	2		19704	N14308		529	
1999997	2003	4	11	14	5		20355	N528AU		1457	
1999998	2012	2	5	15	2		19393	N281WN		536	
1999999	2003	2	4	29	2		19977	N364UA		1241	

1580411 rows × 29 columns

## Data Prep for Model

```
In [15]: #import needed packages  
from sklearn.model_selection import train_test_split  
from sklearn.preprocessing import LabelEncoder
```

```
In [16]: #turns tail number into a numeric representation  
label_encoder = LabelEncoder()  
data_df['tail_number'] = label_encoder.fit_transform(data_df['tail_number'])
```

```
In [17]: #Creates feature and target variables  
x = data_df.drop('arrdelay', axis=1)  
y = data_df['arrdelay']
```

```
In [18]: #Splits the data into training and testing sets  
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)
```

## Model Building

```
In [34]: #import needed packages  
from sklearn.preprocessing import MinMaxScaler  
from sklearn.model_selection import GridSearchCV  
from sklearn.pipeline import Pipeline  
from xgboost import XGBRegressor  
from catboost import CatBoostRegressor  
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

```
In [27]: #Creates a pipeline with different models  
pipeline = Pipeline([("model", XGBRegressor())]) #place holder since none dose not work  
#Defines parameters for the grid search for both models  
params = [{'model': [XGBRegressor()],  
            'model__n_estimators': [100, 500],  
            'model__learning_rate': [0.3, 0.1]},  
          {'model': [CatBoostRegressor()],  
            'model__iterations': [100, 500],  
            'model__learning_rate': [0.3, 0.1]}]
```

```
In [29]: #Performs a grid search for all the models  
grid_search = GridSearchCV(pipeline, param_grid=params, cv=3)  
grid_search.fit(x_train, y_train)
```

```
Out[29]:
GridSearchCV
└─ estimator: Pipeline
   └─ XGBRegressor
```

```
In [30]: #Evaluatse the best parameters and scores
print("Best parameters:", grid_search.best_params_)
print("Best score:", grid_search.best_score_)
```

```
Best parameters: {'model': <catboost.core.CatBoostRegressor object at 0x0000025FBECFD4E0>, 'model__iteration
s': 500, 'model__learning_rate': 0.1}
Best score: 0.385955662554791
```

```
In [31]: #Evaluates on the test set
test_score = grid_search.score(x_test, y_test)
print("Test Score:", test_score)
```

```
Test Score: 0.37458811810437853
```

```
In [32]: #Gets the best model from the grid search
best_model = grid_search.best_estimator_

#Predicts on the train and test sets
y_pred_train = best_model.predict(x_train)
y_pred = best_model.predict(x_test)
```

## Result Interperation

```
In [35]: #Calculates the RMSE, MAE, and R2 on the train set
rmse_train = mean_squared_error(y_train, y_pred_train, squared=False)
mae_train = mean_absolute_error(y_train, y_pred_train)
r2_train = r2_score(y_train, y_pred_train)
#Calculates the RMSE, MAE, and R2 on the test set
rmse_test = mean_squared_error(y_test, y_pred, squared=False)
mae_test = mean_absolute_error(y_test, y_pred)
r2_test = r2_score(y_test, y_pred)
```

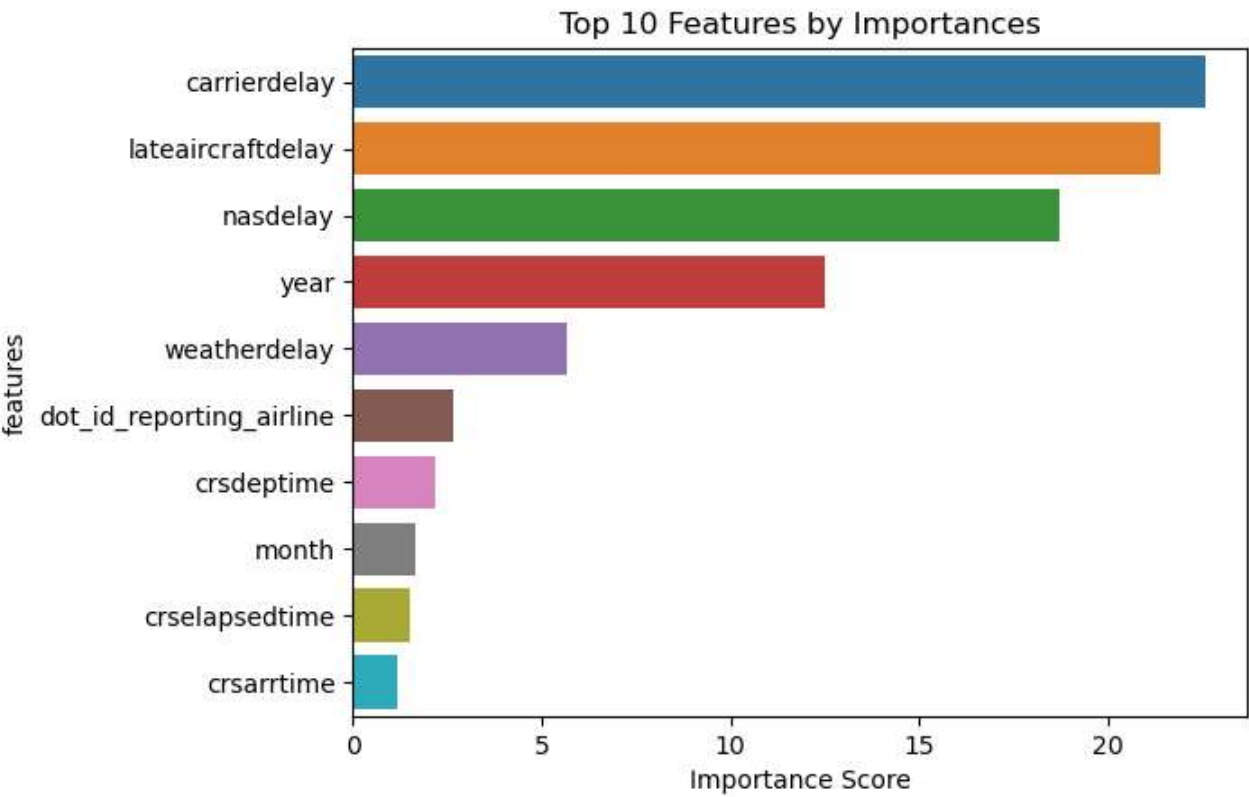
```
In [52]: print("average delay:", data_df['arrdelay'].mean(),"minutes")
print("-----Training data-----")
print("train_RMSE:", rmse_train)
print("train_MAE:", mae_train)
print("train_R2:", r2_train)
print("-----testing data-----")
print("test_RMSE:", rmse_test)
print("test_MAE:", mae_test)
print("test_R2:", r2_test)
```

```
average delay: 6.20160831581152 minutes
-----Training data-----
train_RMSE: 28.363256100141665
train_MAE: 14.45923410486165
train_R2: 0.4049634687084884
-----testing data-----
test_RMSE: 29.997491435380784
test_MAE: 14.621067338119326
test_R2: 0.37458811810437853
```

```
In [56]: #Calculates the feature importance for the model
cat_model = best_model.named_steps['model']
feature_importance = pd.Series(cat_model.feature_importances_, index=x_train.columns)
feature_importance.sort_values(ascending=False, inplace=True)

#Creates a bar plot to display the top 10 features by importances
sn.barplot(x=feature_importance.head(10), y=feature_importance.head(10).index)
plt.xlabel('Importance Score')
plt.ylabel('Features')
plt.title('Top 10 Features by Importances')
```

Out[56]: Text(0.5, 1.0, 'Top 10 Features by Importances')



In [ ]: